

Cyber Threat Hunting Using Large Language Models



Vinayak Tanksale

Abstract Cyber threat hunting plays a critical role in proactively identifying and mitigating potential cybersecurity threats. The emergence of large language models (LLMs) has increased interest in the application of LLMs in cyber threat hunting. This research paper investigates the utilization of LLMs for cyber threat hunting and provides an overview of their potential benefits, challenges, and future directions. LLMs possess powerful language understanding capabilities and can process vast amounts of data, enabling them to uncover hidden patterns, detect anomalies, and identify vulnerabilities. By incorporating LLMs into workflows, organizations can enhance their ability to detect and respond to emerging threats in real time. The adoption of LLMs in threat hunting raises challenges related to bias and fairness, privacy, and computational efficiency. This paper discusses these challenges and explores potential solutions. The paper outlines future directions, such as improving LLM training with cybersecurity-specific domains and integrating contextual knowledge into LLM-based threat hunting approaches.

Keywords Large language models · Cybersecurity · Machine learning · Cyber threat hunting · Artificial intelligence

1 Introduction

The field of cybersecurity is constantly evolving as threat actors devise increasingly sophisticated techniques to bypass traditional security measures. To combat these evolving threats, cybersecurity professionals are continually exploring new approaches and technologies. One such technology that has gained significant attention and holds immense potential is large language models (LLMs). LLMs are advanced artificial intelligence (AI) models trained on vast amounts of text data,

V. Tanksale (✉)
Ball State University, Muncie, IN 47306, USA
e-mail: vjtanksale@bsu.edu
URL: <http://vinayaktanksale.com>

© The Author(s), under exclusive license to Springer Nature Singapore Pte Ltd. 2024
X.-S. Yang et al. (eds.), *Proceedings of Ninth International Congress on Information and Communication Technology*, Lecture Notes in Networks and Systems 1000,
https://doi.org/10.1007/978-981-97-3289-0_50

629

enabling them to understand and generate human-like language. These models have found applications in various domains, including natural language processing, content generation, and now, cyber threat hunting.

Cyber threat hunting is a proactive approach to cybersecurity that focuses on actively searching for and detecting threats that may have bypassed traditional security defenses [2],[?]. It involves analyzing vast amounts of data, identifying anomalies, and investigating potential indicators of compromise (IOCs). Large language models have the potential to enhance cyber threat hunting by augmenting the capabilities of human analysts. LLMs can assist in automating the analysis of security logs, parsing and understanding threat intelligence reports, and identifying patterns or linguistic markers that may indicate malicious activity. By leveraging the power of LLMs, cybersecurity professionals can gain valuable insights, improve the efficiency of their threat hunting efforts, and stay one step ahead of rapidly evolving cyber threats [16].

This research paper aims to explore the intersection of large language models and cyber threat hunting. It investigates the potential applications of LLMs in various stages of the threat hunting lifecycle, including data analysis, anomaly detection, IOC generation, and threat intelligence analysis. Furthermore, it examines the challenges and considerations associated with integrating LLMs into existing cybersecurity frameworks, such as privacy concerns, model biases, and interpretability. By shedding light on the capabilities and limitations of LLMs in the context of cyber threat hunting, this research paper aims to provide valuable insights for cybersecurity professionals seeking to leverage this emerging technology.

By examining the potential of large language models in the context of cyber threat hunting, this research paper contributes to the growing body of knowledge on the application of AI technologies in cybersecurity. Through a comprehensive analysis of the capabilities, limitations, and challenges associated with using LLMs for threat hunting [10][?,?] this research aims to empower cybersecurity professionals with the necessary insights to make informed decisions about integrating LLMs into their existing security frameworks. By harnessing the power of LLMs, cybersecurity professionals can enhance their threat hunting capabilities, improve their defenses against emerging threats, and ultimately safeguard critical assets and data from malicious actors.

2 Overview and Foundations

2.1 *Cyber Threat Hunting*

Cyber threat hunting refers to the proactive and iterative process of detecting, investigating, and mitigating potential cyber threats that evade traditional security defenses [7]. It involves actively searching for signs of malicious activity within a network environment, identifying and analyzing indicators of compromise, and taking neces-

sary actions to neutralize threats. Cyber threat hunting plays a crucial role in modern cybersecurity, as it enables organizations to detect and respond to threats before they can cause significant damage. By taking a proactive approach, cyber threat hunting helps to reduce dwell time, enhance incident response capabilities, and mitigate the potential impact of cyber attacks [8]. It empowers security teams to stay one step ahead of adversaries and protect critical systems and sensitive information. The importance of cyber threat hunting in the contemporary threat landscape cannot be overstated, as it helps organizations to augment their defensive strategies and maintain a resilient cybersecurity posture.

Cyber threat hunting involves a systematic and proactive approach to detect and mitigate potential threats in an organization's network environment [20]. It consists of several key components that collectively contribute to the effectiveness of the hunting process [1]. First, comprehensive threat intelligence plays a crucial role in providing up-to-date information on the latest attack vectors, known threat actors, and emerging vulnerabilities. Second, advanced data analytics techniques, including machine learning algorithms and behavior analytics, help identify anomalous patterns and indicators of compromise within large datasets. Third, robust network and endpoint visibility tools enable the monitoring and analysis of network traffic, system logs, and user activities, providing valuable insights into potential threats. Additionally, collaborative teamwork and effective communication between cybersecurity professionals and IT operations teams facilitate the sharing of knowledge and expertise, enabling a proactive response to threats. The combination of these key components enhances the organization's ability to proactively detect, investigate, and mitigate cyber threats, strengthening its overall cybersecurity posture.

2.2 *Large Language Models*

Large language models [23] are trained using vast amounts of textual data, enabling them to understand and generate human-like text. The training process typically involves utilizing deep learning techniques, such as transformer architectures, which have demonstrated exceptional performance in natural language processing tasks. LLMs like OpenAI's GPT-3 employ a transformer architecture consisting of multiple layers of self-attention mechanisms, enabling them to capture long-range dependencies and contextual information within the text. During training, LLMs are exposed to diverse datasets, including books, articles, and web content, allowing them to learn patterns, grammar, and semantics [5]. The training process incorporates techniques like unsupervised learning and self-supervised learning, where LLMs learn to predict missing words or generate coherent text based on context [19]. The resulting trained LLMs can then be fine-tuned for specific applications, such as cyber threat hunting, to leverage their language processing capabilities and contextual understanding.

LLMs have begun to play a significant role in cybersecurity, offering new possibilities for threat detection, analysis, and response [14]. LLMs excel at natural language processing tasks, enabling them to analyze and understand vast amounts of textual

data, including security logs, threat intelligence reports, and online discussions. By leveraging their language processing capabilities, LLMs can aid in identifying patterns, detecting anomalies, and uncovering potential cyber threats. They can assist in automated threat hunting, generating security alerts, and even providing contextual information to support incident response activities [21]. LLMs can also contribute to the development of more robust security systems, assisting in the creation of advanced intrusion detection and prevention systems. However, ethical considerations, adversarial attacks, and potential biases within LLMs are challenges that need to be carefully addressed to ensure their responsible and effective integration within cybersecurity practices.

3 LLMs for Cyber Threat Hunting

3.1 *Log Analysis and Anomaly Detection*

Large language models have emerged as powerful tools for log analysis and anomaly detection in cybersecurity [22],[?]. Logs generated by various systems and applications contain valuable information about system activities, user behaviors, and potential security incidents. LLMs can process and interpret log data, enabling organizations to extract meaningful insights and detect anomalies that may indicate security breaches or suspicious activities. By leveraging their natural language processing capabilities, LLMs can understand log entries, identify patterns, and recognize deviations from normal behavior [6],[?]. They can learn the contextual relationships between log events, aiding in the detection of complex attack scenarios and zero-day exploits. Furthermore, LLMs can assist in automating the analysis of log data, reducing the manual effort required for anomaly detection and providing security analysts with timely alerts.

LLMs can be trained on diverse log datasets, including system logs, network logs, and application logs, to learn the normal behavior of the environment [25]. By comparing incoming log entries to the learned representations, LLMs can identify deviations and anomalies. This approach offers a proactive way to detect potential threats and security incidents. Moreover, LLMs can be combined with other machine learning techniques, such as clustering algorithms or outlier detection methods, to enhance the accuracy of anomaly detection and reduce false positives.

In addition to log analysis, LLMs can also be utilized for log generation. They can generate synthetic log entries based on learned patterns, which can be used to augment existing log datasets for training anomaly detection models. Synthetic logs can introduce new scenarios and attack vectors, helping to improve the resilience of anomaly detection systems and prepare them for emerging threats.

However, the application of LLMs in log analysis and anomaly detection also poses challenges. The sheer volume and velocity of log data require efficient processing and scaling of LLMs. Additionally, ensuring the interpretability and explain-

ability of LLM-based anomaly detection models is crucial for building trust and facilitating effective decision-making by security analysts. Ongoing research is exploring techniques to address these challenges and optimize the integration of LLMs into log analysis and anomaly detection pipelines.

3.2 Generation of Indicators of Compromise (IOCs)

Large Language Models have shown great potential in the generation of Indicators of Compromise [15],[?], which are crucial in detecting and mitigating cyber threats [13]. IOCs are specific artifacts or patterns that indicate the presence of malicious activities or compromises within a system or network. LLMs can leverage their language generation capabilities to automatically generate IOCs based on various contextual factors, such as known attack techniques, threat intelligence reports, or historical data [18]. By understanding the semantics and syntax of IOCs, LLMs can generate accurate and contextually relevant indicators, aiding in the identification and prevention of cyber threats.

The use of LLMs in IOC generation offers several advantages. Firstly, LLMs can generate a wide range of IOCs, including file hashes, IP addresses, domain names, URLs, and behavioral patterns, providing a comprehensive set of indicators for threat detection. Secondly, LLMs can adapt to evolving threat landscapes by learning from large-scale datasets and incorporating the latest attack techniques and trends. This adaptability enables LLMs to generate IOCs that reflect the current threat landscape and anticipate emerging threats. Additionally, LLM-generated IOCs can assist in automated threat hunting and incident response, enabling faster and more efficient detection and mitigation of security incidents.

To ensure the effectiveness and quality of LLM-generated IOCs, careful training and fine-tuning of the models are essential. LLMs can be trained on diverse datasets, including historical attack data, threat intelligence reports, and security blogs, to learn the patterns and characteristics of IOCs. Fine-tuning the models with domain-specific data further enhances their ability to generate accurate and relevant IOCs for specific industries or organizations [4],[?]. Moreover, continuous model evaluation and validation are crucial to identify any biases or limitations in IOC generation and improve the overall performance and reliability of LLM-based IOC generation systems.

However, challenges exist in the utilization of LLMs for IOC generation. The potential generation of false-positive or false-negative IOCs is a concern that needs to be addressed. Adequate validation mechanisms and human oversight are necessary to verify the generated IOCs before operational deployment. Additionally, privacy and security considerations should be taken into account when using LLMs for IOC generation, as the models may inadvertently reveal sensitive information or inadvertently learn and reproduce malicious patterns.

Despite these challenges, LLMs offer a promising avenue for improving the efficiency and effectiveness of IOC generation in cybersecurity. By harnessing the language generation capabilities of LLMs, organizations can enhance their threat detection capabilities and proactively defend against evolving cyber threats.

3.3 *Threat Intelligence Analysis*

Threat intelligence analysis plays a crucial role in identifying and mitigating cyber threats. With the emergence of large language models (LLMs), such as GPT-3, there is a growing interest in leveraging these models to enhance threat intelligence analysis. LLMs have shown remarkable capabilities in natural language processing, understanding context, and generating human-like text. This opens up new possibilities for extracting valuable insights from vast amounts of unstructured threat intelligence data.

One key area where LLMs can be utilized in threat intelligence analysis is in automating the process of extracting and categorizing threat intelligence information. LLMs can be trained to understand and interpret threat reports, security blogs, forums, and other sources of threat intelligence. They can analyze these texts, identify relevant indicators of compromise (IOCs), extract key information such as malware names, attacker techniques, and affected systems, and categorize the threats based on severity or threat actor attribution. This automation helps in reducing the manual effort required for threat intelligence analysis and enables faster response to emerging threats.

Furthermore, LLMs can assist in identifying patterns and correlations across different threat intelligence sources. By processing a wide range of textual data, LLMs can identify commonalities in attack vectors, TTPs (tactics, techniques, and procedures), and threat actor behaviors. This can enable analysts to detect emerging trends, identify previously unknown threat campaigns, and proactively update defenses. LLMs can also be utilized to identify context-based relationships between threat actors, campaigns, and targeted industries or sectors, providing valuable insights for targeted threat intelligence.

Another potential application of LLMs in threat intelligence analysis is in the generation of high-quality threat reports. LLMs can generate human-like narratives based on the analyzed threat intelligence data, providing detailed and actionable information to security teams and stakeholders. These generated reports can include summaries of threat campaigns, recommended mitigation strategies, and indicators of compromise. By automating the report generation process, LLMs enable analysts to focus on higher-level analysis and decision-making tasks.

Moreover, LLMs can contribute to the improvement of threat intelligence sharing and collaboration. By understanding and interpreting textual descriptions of threat intelligence, LLMs can help in the extraction of relevant information for sharing with trusted partners and communities. LLMs can assist in anonymizing sensitive data, categorizing threat intelligence based on sharing agreements and standards,

and facilitating the exchange of actionable intelligence in a timely manner. This enhances the collective defense against cyber threats and promotes a proactive and collaborative approach to cybersecurity.

Several research studies have explored the potential of LLMs in threat intelligence analysis. For example, Kim et al. [15] proposed a framework for generating indicators of compromise (IOCs) based on language models. They demonstrated the effectiveness of their approach in extracting IOCs from threat reports. Additionally, Garg et al. [12] presented a method that utilizes LLMs for threat intelligence analysis and categorization, achieving promising results in identifying threat types and classifying threat severity.

The utilization of large language models (LLMs) in threat intelligence analysis offers significant advantages in automating the extraction, categorization, and correlation of threat intelligence data. LLMs can enhance the speed and accuracy of threat analysis, improve the generation of threat reports, and facilitate information sharing and collaboration among security stakeholders. While challenges related to model bias, interpretability, and data privacy need to be addressed, the potential benefits of incorporating LLMs into threat intelligence analysis are promising, paving the way for more advanced and efficient cybersecurity practices.

4 Evaluation

4.1 Metrics

Metrics and criteria play a significant role in assessing the capabilities and limitations of LLMs in cyber threat hunting. This section explores key metrics and criteria that can be used for evaluating LLMs in threat hunting scenarios. One important metric for evaluating LLMs is the accuracy of threat detection. This metric measures the LLM's ability to correctly identify and classify threats from various sources of data, such as log files, network traffic, and security alerts. Accuracy can be determined by comparing the LLM's output with ground truth labels provided by domain experts or by analyzing the true positive and false positive rates. Additionally, metrics such as precision, recall, and F1 score can provide insights into the LLM's performance in identifying true threats while minimizing false positives.

Another crucial criterion is the interpretability of LLMs. Threat hunters need to understand how LLMs arrive at their predictions to gain insights into the underlying threat landscape. Interpretability metrics assess the ability of LLMs to provide explanations for their decisions, such as highlighting important features, providing contextual information, or generating human-readable explanations. Metrics like explanation fidelity, clarity, and relevance can be employed to evaluate the interpretability of LLMs.

The scalability and efficiency of LLMs are also important considerations in threat hunting. Metrics related to model size, inference time, and resource utilization can

gauge the scalability and efficiency of LLMs. These metrics help assess the feasibility of deploying LLMs in real-time threat hunting scenarios, especially when dealing with large volumes of data or time-sensitive investigations. Additionally, criteria such as training time and memory requirements can inform decisions regarding the practicality and resource allocation for LLM deployment.

Robustness and resilience are critical criteria for evaluating LLMs in threat hunting. Robustness refers to the LLM's ability to handle various types of threats, including novel or adversarial examples. Metrics like adversarial robustness, generalization performance, and vulnerability analysis can provide insights into the LLM's resilience against attacks or attempts to deceive the model. Evaluating LLMs under different threat scenarios and perturbations helps assess their robustness in real-world threat hunting environments. Lastly, user satisfaction and usability are subjective but important criteria for evaluating LLMs in threat hunting. User feedback, surveys, and usability testing can be conducted to gather insights on the ease of use, user interface design, and overall satisfaction with the LLM-based threat hunting system. These metrics assess the practicality and user-friendliness of LLMs, considering the perspectives and needs of threat hunters and security analysts.

Several research studies have contributed to evaluating LLMs in the context of threat hunting. For instance, Chen et al. [6] proposed a methodology for evaluating the accuracy and interpretability of LLMs in threat detection tasks. They developed metrics based on interpretability fidelity and semantic consistency to assess LLM performance. Additionally, Nguyen et al. [17] presented a framework for evaluating the scalability and efficiency of LLMs in large-scale log analysis, considering factors such as model size, processing time, and memory usage.

Metrics and criteria for evaluating LLMs in threat hunting encompass accuracy, interpretability, scalability, robustness, and user satisfaction. These metrics provide insights into the LLM's performance, explainability, efficiency, resilience, and usability in real-world threat hunting scenarios. By considering these evaluation aspects, organizations can make informed decisions regarding the deployment and utilization of LLMs in their threat hunting processes, ultimately enhancing their cybersecurity posture.

4.2 Real-Time Analysis

The performance and efficiency of large language models in real-time analysis are critical factors that determine their effectiveness in processing and analyzing data in a timely manner. This section explores the performance and efficiency aspects of LLMs in real-time analysis and their implications for various applications. One key aspect of performance is the speed of LLM inference. Real-time analysis requires LLMs to process and generate predictions within strict time constraints. The inference speed of LLMs can be measured in terms of tokens per second or queries per second. Faster inference enables LLMs to handle high-volume data streams and respond to

threats in real time. Techniques such as model optimization, parallel processing, and hardware acceleration can be employed to improve the inference speed of LLMs.

Another important consideration is the memory footprint of LLMs. Efficient memory utilization is crucial for real-time analysis as it directly impacts the scalability and responsiveness of the system. LLMs with large memory requirements may face challenges when processing large-scale datasets or operating in resource-constrained environments. Techniques such as model compression, quantization, and sparsity optimization can help reduce the memory footprint of LLMs without significantly sacrificing performance.

The scalability of LLMs in real-time analysis is a crucial factor, especially in scenarios with rapidly increasing data volumes. Scalability refers to the ability of LLMs to handle larger datasets or higher query rates without compromising performance. Distributed computing frameworks and parallel processing techniques can be employed to scale LLM-based systems horizontally, allowing them to process data in parallel and handle increased workloads effectively.

Energy efficiency is another dimension of efficiency that is gaining importance in real-time analysis. LLMs often require substantial computational resources, which can result in high energy consumption. Optimizing the energy efficiency of LLMs can be achieved through techniques such as model compression, low-power hardware architectures, and intelligent resource allocation. Energy-efficient LLMs enable longer operational times and reduce the environmental impact of real-time analysis systems. Dai et al. [9] proposed a method to improve the inference speed of LLMs by leveraging multi-threading and GPU acceleration. They demonstrated significant speedups in real-time sentiment analysis tasks. Wang et al. [24] introduced a framework for efficient and scalable LLM-based anomaly detection, which utilized parallel computing techniques and achieved high performance in processing large-scale log data.

The performance and efficiency of LLMs in real-time analysis are crucial for their successful deployment in various applications. Improving the inference speed, memory footprint, scalability, and energy efficiency of LLMs enables timely and effective analysis of data streams, enhancing the responsiveness and reliability of real-time analysis systems. By considering these aspects, organizations can select and optimize LLM-based solutions that meet their specific requirements and operational constraints.

4.3 Bias and Fairness

Assessing model bias and fairness in threat detection using large language models is crucial to ensure equitable and unbiased outcomes in cybersecurity systems. LLMs have gained popularity due to their ability to process and analyze vast amounts of textual data, including threat intelligence reports and online content. However, like any machine learning model, LLMs can be susceptible to bias, and it is essential to evaluate their fairness to avoid discriminatory outcomes.

One approach to assessing model bias and fairness in LLMs is to examine the training data. Biases present in the training data can be reflected in the model's outputs. Therefore, it is important to ensure that the training data is diverse, representative, and free from biases. Furthermore, methods such as debiasing techniques and data augmentation can be applied to reduce biases and promote fairness during the training phase of LLMs.

Several metrics and evaluation techniques can be employed to assess model bias and fairness in LLMs. For example, the disparate impact ratio measures the ratio of positive predictions for different demographic groups. The equal opportunity difference quantifies the difference in true positive rates across different groups. Other metrics such as the false positive and false negative rates can also be used to evaluate fairness. Additionally, fairness-aware fine-tuning methods can be applied to mitigate biases and enhance fairness in LLMs.

Research has been conducted on assessing model bias and fairness in LLMs. Bolukbasi et al. [3] proposed a method to learn gender-neutral word embeddings, aiming to mitigate gender biases in natural language processing tasks. Gonen and Goldberg [11] introduced a method for quantifying and mitigating biases in word embeddings. These studies provide insights and techniques for addressing biases in LLMs and promoting fairness in threat detection applications.

To assess model bias and fairness in LLMs, it is crucial to consider the social and ethical implications. Transparency in model development, documentation, and decision-making processes is vital to ensure accountability and enable external auditing. Additionally, involving diverse stakeholders and experts from various domains can help identify and address potential biases in LLMs, fostering a collaborative and inclusive approach toward fairness.

Assessing model bias and fairness in threat detection using LLMs is essential for promoting equitable and unbiased outcomes. By applying appropriate evaluation metrics, considering diverse training data, and employing fairness-aware techniques, biases in LLMs can be identified and mitigated. The continuous efforts toward assessing and enhancing fairness in LLMs contribute to the development of more reliable and unbiased threat detection systems, promoting fairness and equality in cybersecurity.

5 Future Directions and Challenges

One future direction for large language models in cyber threat hunting is the integration of explainability and interpretability techniques. As LLMs become more sophisticated, understanding the reasoning behind their predictions becomes crucial. Research, such as the work by Ribeiro et al. [21], on explaining the predictions of complex models can be leveraged to provide insights into the decision-making process of LLMs and improve their transparency.

Enhancing the robustness and resilience of LLMs to adversarial attacks is an ongoing challenge. Adversarial attacks aim to manipulate LLMs' behavior by intro-

ducing subtle changes to input data. Research, such as the work by Mateusz et al. [14], on poisoning attacks and countermeasures can inform the development of defense mechanisms to mitigate the impact of such attacks on LLMs used in cyber threat hunting.

Scaling LLMs to handle the increasing volume and complexity of cybersecurity data is a pressing future direction. Efficient architectures and distributed computing techniques, such as those explored by Bonawitz et al. [4] in federated learning, can enable LLMs to process and analyze massive datasets while maintaining privacy and security. Addressing ethical considerations and biases in LLMs is crucial for their responsible deployment in cyber threat hunting.

Collaborative efforts and shared datasets are essential for advancing LLMs in cyber threat hunting. Incorporating real-time analysis capabilities into LLMs for proactive threat hunting is an exciting future direction. Real-time monitoring and analysis can enable faster response to emerging threats. Metrics and criteria specific to LLMs can be devised to assess their performance, accuracy, and resilience against attacks.

The use of LLMs in threat hunting raises important privacy concerns that need to be addressed. LLMs are trained on large amounts of data, which may include sensitive and private information. The models have the potential to learn and retain information about individuals and organizations, leading to privacy risks if not handled properly. It is crucial to establish robust privacy protection mechanisms when utilizing LLMs in threat hunting scenarios. To address privacy concerns, researchers and practitioners need to explore methods for privacy-preserving LLM training and deployment.

While LLMs have shown tremendous potential in cyber threat hunting, future research should focus on enhancing interpretability, addressing bias and fairness, improving scalability and efficiency, developing evaluation methodologies, integrating LLMs with other cybersecurity domains, and ensuring ethical use. These advancements will pave the way for more effective and responsible utilization of LLMs in cyber threat hunting. Techniques for data anonymization and differential privacy can be applied to ensure that sensitive information is not directly exposed during the training and deployment of LLMs. It is crucial to strike a balance between the utility of LLMs in threat hunting and the protection of individuals' privacy rights.

6 Conclusion

Large language models have emerged as powerful tools in the field of cyber threat hunting. Their ability to process and analyze vast amounts of text data, combined with their language understanding capabilities, makes them valuable assets in detecting and mitigating cyber threats. LLMs offer several advantages, including their capacity to uncover hidden patterns, identify indicators of compromise, and aid in anomaly detection.

However, the adoption of LLMs in cyber threat hunting also presents challenges and considerations. These include the need for interpretability and explainability of

LLM predictions, addressing model bias and fairness, ensuring privacy protection, and dealing with the trade-off between model performance and computational efficiency. To fully harness the potential of LLMs, researchers and practitioners should continue to explore and develop innovative approaches in these areas.

Moreover, future research should focus on improving the training process of LLMs with respect to cybersecurity-specific domains, as well as integrating domain-specific knowledge and context into their learning. Collaboration between cybersecurity experts, data scientists, and language model developers is crucial to enhance the effectiveness and applicability of LLMs in threat hunting.

Despite the challenges and considerations, LLMs hold significant promise for the future of cyber threat hunting. Their ability to continuously learn and adapt to evolving threats, combined with the advancements in interpretability and privacy-preserving techniques, can greatly enhance organizations' cybersecurity posture. By leveraging LLMs effectively, organizations can gain valuable insights, proactively detect and respond to threats, and ultimately strengthen their defense against cyber-attacks.

The integration of large language models into cyber threat hunting represents an exciting and rapidly evolving field. With continued research, development, and collaboration, LLMs have the potential to revolutionize the way we approach cybersecurity, enabling more proactive and effective threat detection and response.

References

1. Baccas T, Gharakheili Habib H, Armita S (2021) Cyber threat hunting: a systematic review. *Comput Secur* 106:102295
2. Bandyopadhyay D, Joshi A, Dutt R, Chakraborty A (2019) Cyber threat hunting: an overview. *Int J Netw Secur Appl* 11(6):1–20
3. Bolukbasi T, Chang K-W, Zou JY, Saligrama V, Kalai A (2016) Man is to computer programmer as woman is to homemaker? debiasing word embeddings. In: *Advances in neural information processing systems*, pp 4349–4357
4. Bonawitz K, Ivanov V, Kreuter B, Marcedone A, Brendan McMahan H, Patel S, Yarmolinsky A (2019) Towards federated learning at scale: system design. *arXiv preprint* [arXiv:1902.01046](https://arxiv.org/abs/1902.01046)
5. Brown TB, Mann B, Ryder N, Subbiah M, Kaplan J, Dhariwal P, Amodei D (2020) Language models are few-shot learners. *arXiv preprint* [arXiv:2005.14165](https://arxiv.org/abs/2005.14165)
6. Chen Z, Bai Y, Zhou L, Wu Q (2020) Logbert+: pretrained natural language understanding for log anomaly detection. *arXiv preprint* [arXiv:2010.13002](https://arxiv.org/abs/2010.13002)
7. Cole E, Ring M (2014) Hiding in plain sight: cyber threat hunting. *IEEE Secur Privacy* 12(4):74–77
8. Crawford J (2019) Cyber threat hunting and incident response: defending organizations against advanced threats. *IT Professional* 21(5):24–31
9. Dai Z, Yang Z, Yang Y, Carbonell J, Le QV, Salakhutdinov R (2020) Transformer-xl: attentive language models beyond a fixed-length context. In: *Proceedings of the 57th annual meeting of the association for computational linguistics*, pp 2978–2988
10. Garg A, Saha SB, Poria S (2022) Large language models: a review. *ACM Trans Asian Low-Resour Lang Inf Process (TALLIP)* 21(1):1–32
11. Gonen H, Goldberg Y (2019) Lipstick on a pig: debiasing methods cover up systematic gender biases in word embeddings but do not remove them. *arXiv preprint* [arXiv:1903.03862](https://arxiv.org/abs/1903.03862)

12. Gwak J, Park J, Lee H, Kim J (2020) Towards automated threat intelligence analysis using large-scale language models. In: Proceedings of the 2020 IEEE international conference on big data (big data)
13. Islam R, Mahbub S, Ahmed I (2020) Generation of cyber threat intelligence indicators of compromise using LSTM and GAN. *IEEE Access* 8:96888–96898
14. Jagielski M, Oprea A, Biggio B, Liu C, Nita-Rotaru C, Li B (2019) Manipulating machine learning: poisoning attacks and countermeasures for regression learning. *ACM Trans Intell Syst Technol* 10(3):1–27
15. Kim Y, Yun H, Lee Y, Kim J, Kim D (2021) Indicators of compromise generation based on language models. *J Inf Sci Eng* 37(5):1163–1177
16. Montenegro V, Ortiz M et al (2018) Hunting cyber criminals: a hacker’s guide to online intelligence gathering tools and techniques. Syngress
17. Nguyen DC, Tran TT, Phung D (2021) Sc-transformer: scalable log analysis with pretrained language models. *arXiv preprint arXiv:2103.07535*
18. Panda A, Ghosh S, Dasgupta D (2021) Automated indicator generation using transformer based language models for cybersecurity operations. In: 2021 IEEE symposium series on computational intelligence (SSCI)
19. Radford A, Wu J, Child R, Luan D, Amodei D, Sutskever I (2019) Language models are unsupervised multitask learners. *OpenAI Blog* 1(8)
20. Ramachandran P (2017) Cyber threat hunting: a proactive approach to advanced threats. *IEEE Secur Privacy* 15(5):46–53
21. Ribeiro MT, Singh S, Guestrin C (2016) “Why should i trust you?” explaining the predictions of any classifier. In: Proceedings of the 22nd ACM SIGKDD international conference on knowledge discovery and data mining. ACM, pp 1135–1144
22. Sun J, Zhu Y, Zhu W, Liu S, Tan Y (2021) Logbert: transformer-based log anomaly detection. *IEEE Trans Ind Inform* 17(10):6771–6780
23. Vaswani A, Shazeer N, Parmar N, Uszkoreit J, Jones L, Gomez AN, Polosukhin I (2017) Attention is all you need. In: Advances in neural information processing systems, pp 5998–6008 (2017) Attention is all you need. In: Advances in neural information processing systems, pp 5998–6008
24. Wang K, Zheng S, Liu Z, Zhou L, Zhang X (2021) Gluoncv and gluonnlp: deep learning in computer vision and natural language processing. *J Mach Learn Res* 22(137):1–7
25. Yang M, Zheng S, Wang K, Zhang X, Zhou L (2021) Deeplog: anomaly detection and diagnosis from system logs through deep learning. *ACM Trans Internet Technol* 21(4):1–26